

# Classification – Chapter 3 Summary

Hands-On Machine Learning with Scikit-Learn and TensorFlow

Study Notes Summary

June 29, 2025

## Contents

<b>1</b>	<b>The Classification Problem</b>	<b>2</b>
1.1	Types of Classification . . . . .	2
1.2	MNIST Dataset Example . . . . .	2
<b>2</b>	<b>Training a Binary Classifier</b>	<b>2</b>
2.1	SGD Classifier . . . . .	2
2.2	Decision Functions vs Predicted Classes . . . . .	3
<b>3</b>	<b>Performance Measures</b>	<b>3</b>
3.1	Accuracy and Its Limitations . . . . .	3
3.2	Confusion Matrix . . . . .	4
3.3	Precision, Recall, and F1 Score . . . . .	4
3.4	Precision/Recall Tradeoff . . . . .	5
3.5	ROC Curve and AUC . . . . .	5
<b>4</b>	<b>Using Cross-Validation</b>	<b>6</b>
<b>5</b>	<b>Multiclass Classification</b>	<b>7</b>
5.1	Strategies for Multiclass Classification . . . . .	7
<b>6</b>	<b>Error Analysis</b>	<b>8</b>
<b>7</b>	<b>Multilabel and Multioutput Classification</b>	<b>9</b>
7.1	Multilabel Classification . . . . .	9
7.2	Multioutput Classification . . . . .	9
<b>8</b>	<b>Key Tools and Code Summary</b>	<b>10</b>
8.1	Essential Scikit-Learn Tools . . . . .	10
<b>9</b>	<b>Key Takeaways</b>	<b>11</b>
9.1	Chapter Summary . . . . .	11
9.2	Best Practices . . . . .	11
9.3	When to Use Different Metrics . . . . .	12

# 1 The Classification Problem

Classification is a supervised learning task where the goal is to predict discrete categories or classes for given input data.

## 1.1 Types of Classification

- **Binary Classification:** Distinguishing between two classes (e.g., spam vs. not spam)
- **Multiclass Classification:** Distinguishing between more than two classes (e.g., digit recognition 0-9)
- **Multilabel Classification:** Each instance can belong to multiple classes simultaneously
- **Multioutput Classification:** Multiple outputs, each with multiple possible classes

## 1.2 MNIST Dataset Example

The MNIST dataset contains 70,000 images of handwritten digits (0-9), each 28×28 pixels. It's a classic benchmark for classification algorithms.

Listing 1: Loading MNIST Dataset

```
1 from sklearn.datasets import fetch_openml
2 import numpy as np
3
4 # Load MNIST dataset
5 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
6 X, y = mnist["data"], mnist["target"]
7
8 # Convert target to integers
9 y = y.astype(np.uint8)
10
11 # Split into train and test sets
12 X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```

### Tip

The MNIST dataset is already shuffled, but it's good practice to shuffle your data before training to ensure the learning algorithm doesn't get biased by the order of instances.

# 2 Training a Binary Classifier

Let's start with a binary classification problem: detecting whether a digit is a 5 or not.

Listing 2: Creating Binary Classification Target

```
1 # Create binary target: True for 5, False for all other digits
2 y_train_5 = (y_train == 5)
3 y_test_5 = (y_test == 5)
```

## 2.1 SGD Classifier

The Stochastic Gradient Descent (SGD) classifier is efficient for large datasets and can handle various loss functions.

Listing 3: Training SGD Classifier

```

1 from sklearn.linear_model import SGDClassifier
2
3 # Create and train SGD classifier
4 sgd_clf = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42)
5 sgd_clf.fit(X_train, y_train_5)
6
7 # Make predictions
8 some_digit = X[0]
9 prediction = sgd_clf.predict([some_digit])
10 print(f"Prediction: {prediction[0]}") # True or False

```

## 2.2 Decision Functions vs Predicted Classes




Figure 1: Decision Function Threshold

Listing 4: Decision Function vs Prediction

```

1 # Get decision scores
2 decision_scores = sgd_clf.decision_function([some_digit])
3 print(f"Decision score: {decision_scores[0]}")
4
5 # Prediction is based on whether score > 0
6 prediction = sgd_clf.predict([some_digit])
7 print(f"Prediction: {prediction[0]}")

```

## 3 Performance Measures

### 3.1 Accuracy and Its Limitations

Accuracy is the ratio of correct predictions to total predictions, but it can be misleading with imbalanced datasets.

Listing 5: Measuring Accuracy

```

1 from sklearn.metrics import accuracy_score
2 from sklearn.model_selection import cross_val_score
3
4 # Cross-validation accuracy
5 accuracy_scores = cross_val_score(sgd_clf, X_train, y_train_5,
6                                   cv=3, scoring="accuracy")
7 print(f"Accuracy: {accuracy_scores.mean():.4f}")

```

#### Warning

For imbalanced datasets (like our 5-detector where only 10% are 5s), accuracy can be misleading. A classifier that always predicts "not 5" would achieve 90% accuracy!

### 3.2 Confusion Matrix

A confusion matrix provides detailed breakdown of correct and incorrect predictions.

Listing 6: Confusion Matrix

```

1 from sklearn.metrics import confusion_matrix
2 from sklearn.model_selection import cross_val_predict
3
4 # Get predictions using cross-validation
5 y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)
6
7 # Generate confusion matrix
8 cm = confusion_matrix(y_train_5, y_train_pred)
9 print("Confusion Matrix:")
10 print(cm)

```

		Negative	Positive
Actual	Negative	<b>TN</b>	<b>FP</b>
	Positive	<b>FN</b>	<b>TP</b>
		Predicted Negative	Predicted Positive

Figure 2: Confusion Matrix Structure

### 3.3 Precision, Recall, and F1 Score

$$\text{Precision} = \frac{TP}{TP + FP} \quad (1)$$

$$\text{Recall (Sensitivity)} = \frac{TP}{TP + FN} \quad (2)$$

$$\text{F1 Score} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (3)$$

Listing 7: Precision

```

1 from sklearn.metrics import precision_score, recall_score, f1_score
2
3 # Calculate metrics
4 precision = precision_score(y_train_5, y_train_pred)
5 recall = recall_score(y_train_5, y_train_pred)
6 f1 = f1_score(y_train_5, y_train_pred)
7
8 print(f"Precision: {precision:.4f}")
9 print(f"Recall: {recall:.4f}")
10 print(f"F1 Score: {f1:.4f}")

```

### 3.4 Precision/Recall Tradeoff

There's always a tradeoff between precision and recall. Adjusting the decision threshold affects both metrics.




Figure 3: Precision/Recall Tradeoff

Listing 8: Adjusting Decision Threshold

```

1 from sklearn.metrics import precision_recall_curve
2
3 # Get decision scores for all training instances
4 y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,
5                             method="decision_function")
6
7 # Calculate precision and recall for different thresholds
8 precisions, recalls, thresholds = precision_recall_curve(y_train_5, y_scores)
9
10 # Find threshold for 90% precision
11 threshold_90_precision = thresholds[np.argmax(precisions >= 0.90)]
12 print(f"Threshold for 90% precision: {threshold_90_precision}")
13
14 # Make predictions with custom threshold
15 y_train_pred_90 = (y_scores >= threshold_90_precision)
16 precision_90 = precision_score(y_train_5, y_train_pred_90)
17 recall_90 = recall_score(y_train_5, y_train_pred_90)
18 print(f"Precision: {precision_90:.4f}, Recall: {recall_90:.4f}")

```

### 3.5 ROC Curve and AUC

The Receiver Operating Characteristic (ROC) curve plots True Positive Rate vs False Positive Rate.

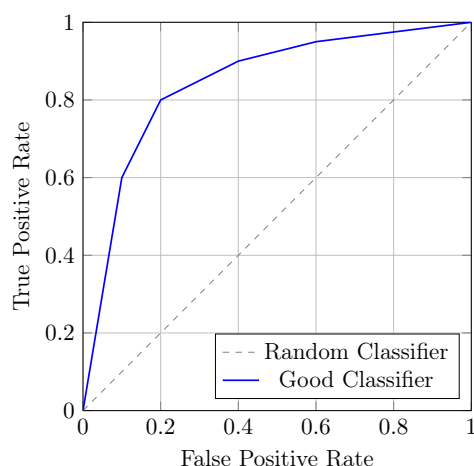


Figure 4: ROC Curve Example

Listing 9: ROC Curve and AUC

```

1 from sklearn.metrics import roc_curve, roc_auc_score
2
3 # Calculate ROC curve
4 fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)
5
6 # Calculate AUC
7 auc_score = roc_auc_score(y_train_5, y_scores)
8 print(f"AUC Score: {auc_score:.4f}")

```

**Note**

Use ROC/AUC when classes are roughly balanced. Use Precision/Recall curves when dealing with imbalanced datasets or when false positives are more important than false negatives.

## 4 Using Cross-Validation

Cross-validation provides more robust performance estimates by training and testing on different data splits.

Listing 10: Cross-Validation for Classification

```

1 from sklearn.model_selection import StratifiedKFold
2
3 # Stratified K-Fold ensures balanced class distribution in each fold
4 skfolds = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
5
6 for train_index, test_index in skfolds.split(X_train, y_train_5):
7     clone_clf = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42)
8     X_train_folds = X_train[train_index]
9     y_train_folds = y_train_5[train_index]
10    X_test_fold = X_train[test_index]
11    y_test_fold = y_train_5[test_index]
12
13    clone_clf.fit(X_train_folds, y_train_folds)
14    y_pred = clone_clf.predict(X_test_fold)
15    n_correct = sum(y_pred == y_test_fold)
16    print(f"Accuracy: {n_correct/len(y_pred):.4f}")

```

Listing 11: Using `cross_val_predict` for *ConfusionMatrix*

```

1 # Get cross-validated predictions
2 y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)
3
4 # Calculate confusion matrix
5 cm = confusion_matrix(y_train_5, y_train_pred)
6 print("Confusion Matrix:")
7 print(cm)
8
9 # Generate classification report
10 from sklearn.metrics import classification_report
11 print("\nClassification Report:")
12 print(classification_report(y_train_5, y_train_pred))

```

## 5 Multiclass Classification

### 5.1 Strategies for Multiclass Classification

- **One-vs-Rest (OvR)**: Train one binary classifier per class
- **One-vs-One (OvO)**: Train one binary classifier for each pair of classes

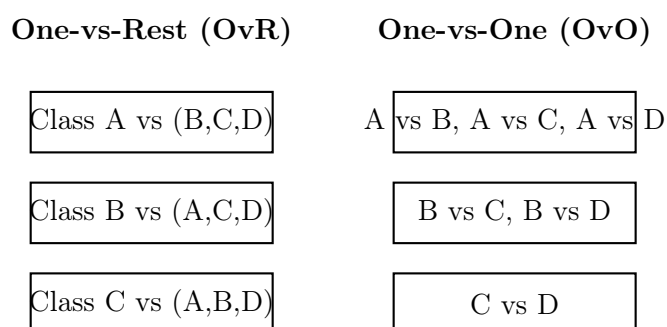


Figure 5: Multiclass Classification Strategies

Listing 12: Multiclass Classification with SGD

```

1 # Train multiclass classifier (automatically uses OvR for SGD)
2 sgd_clf_multi = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42)
3 sgd_clf_multi.fit(X_train, y_train)
4
5 # Make prediction
6 some_digit_pred = sgd_clf_multi.predict([some_digit])
7 print(f"Predicted digit: {some_digit_pred[0]}")
8
9 # Get decision scores for all classes
10 decision_scores = sgd_clf_multi.decision_function([some_digit])
11 print(f"Decision scores: {decision_scores}")
12 print(f"Predicted class: {np.argmax(decision_scores)}")

```

Listing 13: Forcing OvO Strategy

```

1 from sklearn.multiclass import OneVsOneClassifier
2
3 # Force OvO strategy
4 ovo_clf = OneVsOneClassifier(SGDClassifier(max_iter=1000, tol=1e-3,
5     random_state=42))

```

```

5 ovo_clf.fit(X_train, y_train)
6 prediction = ovo_clf.predict([some_digit])
7 print(f"OvO Prediction: {prediction[0]}")
8 print(f"Number of binary classifiers: {len(ovo_clf.estimators_)}")

```

### Tip

Scikit-Learn automatically detects when you use a binary classifier for multiclass tasks and uses the appropriate strategy (OvR for most classifiers, OvO for SVM).

## 6 Error Analysis

Error analysis helps identify patterns in classification mistakes and guide improvements.

Listing 14: Multiclass Confusion Matrix

```

1 # Get cross-validated predictions for multiclass
2 y_train_pred_multi = cross_val_predict(sgd_clf_multi, X_train, y_train, cv=3)
3
4 # Generate confusion matrix
5 cm_multi = confusion_matrix(y_train, y_train_pred_multi)
6 print("Multiclass Confusion Matrix:")
7 print(cm_multi)

```




Figure 6: MNIST Confusion Matrix Visualization

Listing 15: Error Analysis - Normalized Confusion Matrix

```

1 import matplotlib.pyplot as plt
2
3 # Normalize confusion matrix by row (true class)
4 row_sums = cm_multi.sum(axis=1, keepdims=True)
5 norm_cm = cm_multi / row_sums
6
7 # Fill diagonal with zeros to focus on errors
8 np.fill_diagonal(norm_cm, 0)
9
10 # Plot the error-focused confusion matrix
11 plt.figure(figsize=(8,8))
12 plt.matshow(norm_cm, cmap=plt.cm.gray)

```



```

13 plt.colorbar()
14 plt.title("Error Analysis - Normalized Confusion Matrix")
15 plt.xlabel("Predicted")
16 plt.ylabel("Actual")
17 plt.show()

```

### Note

Look for patterns in the confusion matrix:

- Which classes are most often confused?
- Are there systematic errors (e.g., 8s mistaken for 3s)?
- Can you improve preprocessing or feature engineering?

## 7 Multilabel and Multioutput Classification

### 7.1 Multilabel Classification

In multilabel classification, each instance can belong to multiple classes simultaneously.

Listing 16: Multilabel Classification Example

```

1 from sklearn.neighbors import KNeighborsClassifier
2 import numpy as np
3
4 # Create multilabel targets (large digits >= 7, odd digits)
5 y_train_large = (y_train >= 7)
6 y_train_odd = (y_train % 2 == 1)
7 y_multilabel = np.c_[y_train_large, y_train_odd]
8
9 # Train KNN classifier for multilabel task
10 knn_clf = KNeighborsClassifier()
11 knn_clf.fit(X_train, y_multilabel)
12
13 # Make prediction
14 prediction = knn_clf.predict([some_digit])
15 print(f"Multilabel prediction: {prediction}")
16 # Output: [[False, True]] means not large (< 7) but odd

```

### 7.2 Multioutput Classification

Multioutput classification is a generalization where each output can have multiple classes.

Listing 17: Multioutput Classification - Noise Removal

```

1 # Add noise to MNIST images
2 noise = np.random.randint(0, 100, (len(X_train), 784))
3 X_train_mod = X_train + noise
4 noise = np.random.randint(0, 100, (len(X_test), 784))
5 X_test_mod = X_test + noise
6
7 # Target is the original clean image
8 y_train_mod = X_train
9 y_test_mod = X_test
10
11 # Train classifier to remove noise
12 knn_clf.fit(X_train_mod, y_train_mod)
13

```

```

14 # Clean a noisy image
15 clean_digit = knn_clf.predict([X_test_mod[0]])

```

## 8 Key Tools and Code Summary

### 8.1 Essential Scikit-Learn Tools

Tool	Purpose
SGDClassifier	Stochastic Gradient Descent classifier
cross_val_score	Cross-validation scoring
cross_val_predict	Cross-validation predictions
confusion_matrix	Generate confusion matrix
classification_report	Comprehensive classification metrics
precision_score	Calculate precision
recall_score	Calculate recall
f1_score	Calculate F1 score
precision_recall_curve	Precision-recall curve data
roc_curve	ROC curve data
roc_auc_score	Area under ROC curve

Table 1: Key Classification Tools

Listing 18: Complete Classification Pipeline

```

1 from sklearn.linear_model import SGDClassifier
2 from sklearn.model_selection import cross_val_score, cross_val_predict
3 from sklearn.metrics import (confusion_matrix, classification_report,
4                             precision_score, recall_score, f1_score,
5                             roc_auc_score, precision_recall_curve)
6
7 # 1. Train classifier
8 clf = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42)
9 clf.fit(X_train, y_train)
10
11 # 2. Cross-validation evaluation
12 cv_scores = cross_val_score(clf, X_train, y_train, cv=3, scoring='accuracy')
13 print(f"CV Accuracy: {cv_scores.mean():.4f} (+/- {cv_scores.std() * 2:.4f})")
14
15 # 3. Get predictions and detailed metrics
16 y_pred = cross_val_predict(clf, X_train, y_train, cv=3)
17 print("\nConfusion Matrix:")
18 print(confusion_matrix(y_train, y_pred))
19 print("\nClassification Report:")
20 print(classification_report(y_train, y_pred))
21
22 # 4. For binary classification, get additional metrics
23 if len(np.unique(y_train)) == 2:
24     y_scores = cross_val_predict(clf, X_train, y_train, cv=3,
25                               method="decision_function")
26     auc = roc_auc_score(y_train, y_scores)
27     print(f"\nAUC Score: {auc:.4f}")
28
29 # 5. Test set evaluation (final step)
30 test_accuracy = clf.score(X_test, y_test)
31 print(f"Test Accuracy: {test_accuracy:.4f}")

```

## 9 Key Takeaways

### 9.1 Chapter Summary

- **Classification Types:** Binary, multiclass, multilabel, and multioutput classification serve different purposes and require different approaches.
- **Performance Metrics:** Accuracy can be misleading with imbalanced data. Use precision, recall, F1-score, and confusion matrices for better evaluation.
- **Precision vs Recall:** There's always a tradeoff. Adjust decision thresholds based on your specific needs (false positives vs false negatives).
- **Cross-Validation:** Essential for robust performance estimation. Use stratified K-fold for classification to maintain class balance.
- **Multiclass Strategies:** Scikit-Learn automatically handles multiclass classification using OvR or OvO strategies.
- **Error Analysis:** Examine confusion matrices to identify patterns in misclassifications and guide improvements.

### 9.2 Best Practices

#### Tip

##### Classification Best Practices:

- Always use cross-validation for model evaluation
- Choose metrics appropriate for your problem (balanced vs imbalanced data)
- Examine confusion matrices to understand model behavior
- Consider the cost of different types of errors in your domain
- Use stratified sampling to maintain class distributions
- Start with simple models (like SGD) before moving to complex ones
- Always evaluate on a held-out test set for final performance assessment

### 9.3 When to Use Different Metrics

Metric	When to Use
Accuracy	Balanced datasets, equal cost for all errors
Precision	When false positives are costly (e.g., spam detection)
Recall	When false negatives are costly (e.g., medical diagnosis)
F1 Score	When you need a balance between precision and recall
ROC/AUC	Balanced datasets, when you need threshold-independent metric
PR Curve	Imbalanced datasets, when positive class is more important

Table 2: Choosing the Right Metric

#### Note

Remember: The choice of evaluation metric should always align with your business objectives and the costs associated with different types of errors in your specific domain.